

# CSSS508, Lecture 7

## Functions

Michael Pearce

(based on slides from Chuck Lanfear)

November 10, 2022



# Reminders

- Homework 6 is due at *midnight* tonight!
- Lab/Office Hours for next week are *modified*:
  - Office Hours: Monday (11/14) 8:30-10:30am; Thursday (11/17) 10am-12pm
  - Lab (11/15): Replaced with additional office hours (above!)
- Homework 7 will be due on on 11/17 at *midnight*
  - You will have time in class today to work on this homework!
  - *No assignment due Thanksgiving!*

# Topics

Last time, we learned about,

1. Why we use loops
2. `for()` loops
3. `while()` loops

Today, we will cover,

1. Aside: Visualizing the Goal
2. Building blocks of functions
3. Simple functions
4. Using functions with `apply()`

# 1. Visualizing the Goal

# Visualizing the Goal

Before you can write effective code, you need to know *exactly* what you want:

- **Goal:** Do I want a single value? vector? one observation per person? per year?
- **Current State:** What do I currently have? matrix, vector? long or wide format?
- **Translate:** How can I take what I have and turn it into my goal?
  - Sketch out the steps!
  - Break it down into little operations

**As we become more advanced coders, this concept is key!!**

**Remember:** *When you're stuck, try searching your problem on Google!!*

## 2. Building blocks of functions

# Why Functions?

R (as well as mathematics in general) is full of functions!

We use functions to:

- Compute summary statistics (`mean()`, `sd()`, `min()`)
- Fit models to data (`lm(Fertility~Agriculture,data=swiss)`)
- Load data (`read_csv()`)
- Create ggplots (`ggplot()`)
- And so much more!!

# Examples of Existing Functions

- `mean()`:
  - Input: a vector
  - Output: a single number
- `dplyr::filter()`:
  - Input: a data frame, logical conditions
  - Output: a data frame with rows removed using those conditions
- `readr::read_csv()`:
  - Input: a file path, optionally variable names or types
  - Output: a data frame containing info read in from file

Each function requires **inputs**, and returns **outputs**



# Why Write Your Own Functions?

Functions encapsulate actions you might perform often, such as:

- Given a vector, compute some special summary stats
- Given a vector and definition of "invalid" values, replace with NA
- Defining a new logical operator

Advanced function applications (not covered in this class):

- Parallel processing
- Generating *other* functions
- Making custom packages containing your functions

# Anatomy of a Function

```
NAME <- function(ARGUMENT1, ARGUMENT2=DEFAULT){  
  BODY  
  return(OUTPUT)  
}
```

- **Name:** What you call the function so you can use it later
- **Arguments** (aka inputs, parameters): things the user passes to the function that affect how it works
  - e.g. ARGUMENT1, ARGUMENT2
  - ARGUMENT2=DEFAULT is example of setting a default value
  - In this example, ARGUMENT1, ARGUMENT2 values won't exist outside of the function
- **Body:** The actual operations inside the function.
- **Output:** The object inside return(). Could be anything (or nothing!)
  - If unspecified, will be the last thing calculated

# 3. Simple functions

# Example 1: Doubling A Number

```
double_x <- function(x){  
  double_x <- x * 2  
  return(double_x)  
}
```

Let's run it!

```
double_x(5)
```

```
## [1] 10
```

```
double_x(NA)
```

```
## [1] NA
```

```
double_x(1:2)
```

```
## [1] 2 4
```

## Example 2: Extract First/Last

```
first_and_last <- function(x) {  
  first <- x[1]  
  last  <- x[length(x)]  
  return(c("first" = first, "last" = last))  
}
```

Test it out:

```
first_and_last(c(4, 3, 1, 8))
```

```
## first last  
##      4    8
```

# Example 2: Testing

## `first_and_last`

What if I give `first_and_last()` a vector of length 1?

```
first_and_last(7)
```

```
## first last  
##      7      7
```

Of length 0?

```
first_and_last(numeric(0))
```

```
## first  
##      NA
```

Maybe we want it to be a little smarter.

# Example 3: Checking Inputs

Let's make sure we get an error message when the vector is too small:

```
smarter_first_and_last <- function(x) {  
  if(length(x) < 2){  
    stop("Input is not long enough!")  
  } else{  
    first <- x[1]  
    last  <- x[length(x)]  
    return(c("first" = first, "last" = last))  
  }  
}
```

`stop()` ceases running the function and prints the text inside as an error message.

# Example 3: Testing Smarter Function

```
smarter_first_and_last(NA)
```

```
## Error in smarter_first_and_last(NA): Input is not long enough!
```

```
smarter_first_and_last(c(4, 3, 1, 8))
```

```
## first  last  
##      4      8
```



# Cracking Open Functions

If you type a function name without any parentheses or arguments, you can see its contents:

```
smarter_first_and_last
```

```
## function(x) {  
##     if(length(x) < 2){  
##         stop("Input is not long enough!")  
##     } else{  
##         first <- x[1]  
##         last  <- x[length(x)]  
##         return(c("first" = first, "last" = last))  
##     }  
## }  
## <bytecode: 0x7f8dd4d33f48>
```

# Activity: Writing A Function

In Olympic diving, a panel of 7 judges provide scores. After removing the worst and best scores, the mean of the remaining scores is given to the diver. We'll write code to calculate this score!

1. Suppose I give you an ordered vector (lowest to greatest) of length 7. How can you keep elements 2 through 6 (i.e., remove the first and last values).
2. Write a function to calculate a diver's score:
  - Input: Ordered vector of length 7
  - Output: Mean score after removing the first and last scores.
  - Checks: Check that the vector has length 7.
3. Calculate the diver's score given `x <- c(1,3:7,10)`

# Activity: My Solution

1. Extract elements 2 through 6:

- **Answer:** Given vector `x`, use `x[2:6]`

2. Function

```
divers_score <- function(x){  
  if(length(x) != 7){  
    stop("x is not of length 7!")  
  } else{  
    x_nofirst_nolast <- x[2:6]  
    return(mean(x_nofirst_nolast))  
  }  
}
```

1. Calculate the diver's score given `x <- c(1,3:7,10)`

```
divers_score(x = c(1,3:7,10))
```

```
## [1] 5
```

## 4. Using functions with `apply()`

# Applying Functions Multiple Times?

Last week, we saw an example where we wanted to take the mean of each column in the `swiss` data:

```
for(col_index in seq_along(swiss)){  
  mean_swiss_col <- mean(swiss[,col_index])  
  names_swiss_col <- names(swiss)[col_index]  
  print(c(names_swiss_col,round(mean_swiss_col,3)))  
}
```

```
## [1] "Fertility" "70.143"  
## [1] "Agriculture" "50.66"  
## [1] "Examination" "16.489"  
## [1] "Education" "10.979"  
## [1] "Catholic" "41.144"  
## [1] "Infant.Mortality" "19.943"
```

*Isn't this kind of complex?!*

# `apply()`, don't loop!

Writing loops can be challenging and prone to bugs!!

The `apply()` can solve this issue:

- **apply** a function to values in each row or column of a matrix
- Doesn't require preallocation
- Can take built-in functions or user-created functions.

# Structure of `apply()`

`apply()` takes 3 arguments:

1. Data (a matrix or data frame)
2. Margin (1 applies function to each *row*, 2 applies to each *column*)
3. Function

```
apply(DATA, MARGIN, FUNCTION)
```

For example,

```
apply(swiss, 2, mean)
```

##	Fertility	Agriculture	Examination	Education
##	70.14255	50.65957	16.48936	10.97872
##	Catholic	Infant.Mortality		
##	41.14383	19.94255		

# Example 1

```
row_max <- apply(swiss,1,max) #maximum in each row
head(row_max,20)
```

##	Courtelary	Delemont	Franches-Mnt	Moutier	Neuveville
##	80.20	84.84	93.40	85.80	76.90
##	Porrentruy	Broye	Glane	Gruyere	Sarine
##	90.57	92.85	97.16	97.67	91.38
##	Veveyse	Aigle	Aubonne	Avenches	Cossonay
##	98.61	64.10	67.50	68.90	69.30
##	Echallens	Grandson	Lausanne	La Vallee	Lavaux
##	72.60	71.70	55.70	54.30	73.00



# Example 2

```
apply(swiss,2,summary) # summary of each column
```

```
##           Fertility Agriculture Examination Education Catholic
## Min.      35.00000      1.20000      3.00000      1.00000      2.15000
## 1st Qu.   64.70000      35.90000     12.00000      6.00000      5.19500
## Median   70.40000      54.10000     16.00000      8.00000     15.14000
## Mean     70.14255      50.65957     16.48936     10.97872     41.14383
## 3rd Qu.   78.45000      67.65000     22.00000     12.00000     93.12500
## Max.     92.50000      89.70000     37.00000     53.00000    100.00000
##
## Infant.Mortality
## Min.           10.80000
## 1st Qu.         18.15000
## Median          20.00000
## Mean            19.94255
## 3rd Qu.         21.70000
## Max.            26.60000
```

*\*Note:* Matrix output!

# Example 3: User-Created Function

```
scores <- matrix(1:21,nrow=3)  
print(scores)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7]  
## [1,]    1    4    7   10   13   16   19  
## [2,]    2    5    8   11   14   17   20  
## [3,]    3    6    9   12   15   18   21
```

```
apply(scores,1,divers_score)
```

```
## [1] 10 11 12
```

# Activity

*These are homework questions!!*

1. What's similar/different between a `for` loop and `apply`?
2. Write an `apply()` function to take the median value of each column in the `cars` dataset
3. Using `ggplot`, make a scatterplot of the `speed` and `dist` variables in `cars`. Then, add an appropriate horizontal and vertical line symbolizing the median value of each variable.

*Hint:* Using the layers `geom_vline(xintercept = )` and `geom_hline(yintercept = )`

# My Answers

1. What's similar/different between a `for` loop and `apply`?

- **Possible Answer:** `for` loops and `apply` can usually accomplish the same tasks: Each uses code iteratively.
- `for` loops are more flexible, in that you can provide different data using indices to the loop; they are harder to write.
- `apply` loops are simpler and don't require preallocation; they require your data be a matrix or dataframe.

2. Write an `apply()` function to take the median value of each column in the `cars` dataset


```
median_cars <- apply(cars,2,median)
median_cars
```

```
## speed  dist
##    15    36
```

# My Answers

## 1. Make a ggplot

```
library(ggplot2)
ggplot(cars, aes(speed, dist)) + geom_point() +
  geom_vline(xintercept = median_cars[1]) +
  geom_hline(yintercept = median_cars[2])
```



# Homework

Time to work on Homework 7!