

# CSSS508, Lecture 8

## Working with Text Data

Michael Pearce

(based on slides from Chuck Lanfear)

May 17, 2022



# Topics

Last time, we learned about,

1. Aside: Visualizing the Goal
2. Building blocks of functions
3. Simple functions
4. Using functions with `apply()`

Today, we will cover,

1. Basics of Strings
2. Strings in Base R
3. Strings in `stringr` (tidyverse)

# 1. Basics of Strings

# Basics of Strings

- A general programming term for a unit of character data is a **string**
  - Strings are a *sequence of characters*
  - In R, "strings" and "character data" are mostly interchangeable.
  - Some languages have more precise distinctions, but we won't worry about that here!
- We can create strings by surrounding text, numbers, spaces, or symbols with quotes!
  - Examples: `"Hello! My name is Michael"` or `"%*$#01234"`

# Basics of Strings

R can treat strings in funny ways!

```
"01" == "1"
```

```
## [1] FALSE
```

```
"01" == 1
```

```
## [1] FALSE
```

```
"1" == 1
```

```
## [1] TRUE
```

*Reminder:* We can check **data types** using the `class()` function!

```
c(class("1"),class(1))
```

```
## [1] "character" "numeric"
```

## 2. Strings in Base R

- `nchar()`
- `substr()`
- `paste()`

# Data: King County Restaurant Inspections!

Today we'll study real data on **food safety inspections in King County**, collected from [data.kingcounty.gov](https://data.kingcounty.gov).

Note these data are *fairly large*. The following code can be used to download the data directly from my Github page:

```
load(url("https://pearce790.github.io/CSS508/Lectures/Lecture8/resta
```

# Quick Examination of the Data

```
names(restaurants)
```

```
## [1] "Name" "Program_Identifier"  
## [3] "Inspection_Date" "Description"  
## [5] "Address" "City"  
## [7] "Zip_Code" "Phone"  
## [9] "Longitude" "Latitude"  
## [11] "Inspection_Business_Name" "Inspection_Type"  
## [13] "Inspection_Score" "Inspection_Result"  
## [15] "Inspection_Closed_Business" "Violation_Type"  
## [17] "Violation_Description" "Violation_Points"  
## [19] "Business_ID" "Inspection_Serial_Num"  
## [21] "Violation_Record_ID" "Grade"  
## [23] "Date"
```

```
dim(restaurants)
```

```
## [1] 258630 23
```



# Quick Examination of the Data

## Good Questions to Ask:

- What does each row represent?
- Is the data in long or wide format?
- What are the key variables?
- How are the data stored? (*data type*)

# nchar()

The `nchar()` function calculates the *number of characters* in a given string.

- `length()` doesn't work with strings!!
- Why not?

```
nchar("Mike Pearce")
```

```
## [1] 11
```

In our `restaurants` data, let's see how many characters are in each zip code:

```
length_zip <- nchar(restaurants$Zip_Code)
table(length_zip)
```

```
## length_zip
```

```
##      5      10
```

```
## 258629      1
```

# substr()

The `substr()` function allows us to extract characters from a string.

For example, we can extract the third through fifth elements of a string as follows:

```
substr("98126",3,5)
```

```
## [1] "126"
```

# substr()

Let's extract the first five characters from each zip code in the restaurants data, and add it to our dataset.

```
library(dplyr)
restaurants$ZIP_5 <- substr(restaurants$Zip_Code,1,5)
restaurants %>% distinct(ZIP_5) %>% head()
```

```
## # A tibble: 6 × 1
##   ZIP_5
##   <chr>
## 1 98126
## 2 98109
## 3 98101
## 4 98032
## 5 98102
## 6 98004
```

# paste()

We combine strings together using `paste()`. By default, it puts a space between different strings.

For example, we can combine `"Michael"` and `"Pearce"` as follows:

```
paste("Michael", "Pearce")
```

```
## [1] "Michael Pearce"
```

# More complex `paste()` commands

There are two additional common arguments to use with `paste()`:

1. `sep=` controls what separates vectors, entry-wise
2. `collapse=` controls if/how multiple outputs are collapsed into a single string.

```
paste("CSSS", "508", sep= "_")
```

```
## [1] "CSSS_508"
```

```
paste(c("CSSS", "STAT"), "508", sep= "_")
```

```
## [1] "CSSS_508" "STAT_508"
```

```
paste(c("CSSS", "STAT"), "508", sep= "_", collapse= " , ")
```

```
## [1] "CSSS_508 , STAT_508"
```

*When do we get one string as output vs. two?*

# paste()

Let's use `paste()` to create complete mailing addresses for each restaurant:

```
restaurants$mailing_address <-  
  paste(restaurants$Address, ", ",  
        restaurants$City, ", WA ", restaurants$ZIP_5,  
        sep = " ")  
restaurants %>% distinct(mailing_address) %>% head()
```

```
## # A tibble: 6 × 1  
##   mailing_address  
##   <chr>  
## 1 2920 SW AVALON WAY, Seattle, WA 98126  
## 2 10 MERCER ST, Seattle, WA 98109  
## 3 1001 FAIRVIEW AVE N Unit 1700A, SEATTLE, WA 98109  
## 4 1225 1ST AVE, SEATTLE, WA 98101  
## 5 18114 E VALLEY HWY, KENT, WA 98032  
## 6 121 11TH AVE E, SEATTLE, WA 98102
```

### 3. Strings in `stringr`

- `str_length()`
- `str_sub()`
- `str_c()`
- `str_to_upper()`, `str_to_lower()`, and `str_to_title()`
- `str_trim()`
- `str_detect()`
- `str_replace()`



# stringr

`stringr` is yet another R package from the Tidyverse (like `ggplot2`, `dplyr`, `tidyr`, `lubridate`, `readr`).

It provides TONS of functions for working with strings:

- Some are equivalent/better versions of Base R functions
- Some can do *fancier* tricks with strings

*Most* `stringr` functions begin with "`str_`" to make RStudio auto-complete more useful.

We'll cover the basics today, but know there's much more out there!

```
library(stringr)
```

# Equivalencies: `str_length()`

`str_length()` is equivalent to `nchar()`:

```
nchar("weasels")
```

```
## [1] 7
```

```
str_length("weasels")
```

```
## [1] 7
```

# Equivalencies: `str_sub()`

`str_sub()` is like `substr()`:

```
str_sub("Washington", 2, 4)
```

```
## [1] "ash"
```

`str_sub()` also lets you put in negative values to count backwards from the end (-1 is the end, -3 is third from end):

```
str_sub("Washington", 4, -3)
```

```
## [1] "hingt"
```

# Equivalencies: `str_c()`

`str_c()` ("string combine") is just like `paste()` but where the default is `sep = ""` (no space!)

```
str_c(c("CSSS", "STAT"), 508)
```

```
## [1] "CSSS508" "STAT508"
```

```
str_c(c("CSSS", "STAT"), 508, sep=" ")
```

```
## [1] "CSSS 508" "STAT 508"
```

```
str_c(c("CSSS", "STAT"), 508, sep = " ", collapse = ", ")
```

```
## [1] "CSSS 508, STAT 508"
```

# Changing Cases

`str_to_upper()`, `str_to_lower()`, `str_to_title()` convert cases, which is often a good idea to do before searching for values:

```
unique_cities <- unique(restaurants$City)
unique_cities %>% head()
```

```
## [1] "Seattle" "SEATTLE" "KENT" "BELLEVUE" "KENMORE" "Issaquah"
```

```
str_to_upper(unique_cities) %>% head()
```

```
## [1] "SEATTLE" "SEATTLE" "KENT" "BELLEVUE" "KENMORE" "ISSAQUAH"
```

```
str_to_lower(unique_cities) %>% head()
```

```
## [1] "seattle" "seattle" "kent" "bellevue" "kenmore" "issaquah"
```

```
str_to_title(unique_cities) %>% head()
```

```
## [1] "Seattle" "Seattle" "Kent" "Bellevue" "Kenmore" "Issaquah"
```

# Whitespace: `str_trim()`

Extra leading or trailing whitespace is common in text data:

```
unique_names <- unique(restaurants$Name)
unique_names %>% head(3)
```

```
## [1] "@ THE SHACK, LLC "      "10 MERCER RESTAURANT"
## [3] "100 LB CLAM"
```

We can remove the whitespace using `str_trim()`:

```
str_trim(unique_names) %>% head(3)
```

```
## [1] "@ THE SHACK, LLC"      "10 MERCER RESTAURANT"
## [3] "100 LB CLAM"
```

# Patterns!

It's common to want to see if a string satisfies a certain *pattern*.

We did this with numeric values earlier in this course!

```
cars %>% filter(speed < 5 | speed > 24)
```

```
##    speed dist
## 1      4     2
## 2      4    10
## 3     25    85
```

```
cars %>% filter(dist > 2 & dist <= 10)
```

```
##    speed dist
## 1      4    10
## 2      7     4
## 3      9    10
```

# Patterns: `str_detect()`

We can do similar pattern-checking using `str_detect()`:

```
str_detect(string, pattern)
```

- `string` is the character string (or vector of strings) we want to examine
- `pattern` is the pattern that we're checking for inside `string`
- Output: TRUE/FALSE vector indicating if pattern was found

```
str_detect(string = c("Hello", "my name", "is Michael"),  
           pattern = "m")
```

```
## [1] FALSE TRUE FALSE
```

```
str_detect(string = c("Hello", "my name", "is Michael"),  
           pattern = "M")
```

```
## [1] FALSE FALSE TRUE
```

Results are case-sensitive!!



# Patterns: `str_detect()`

Let's see which phone numbers are in the 206 area code:

```
unique_phones <- unique(restaurants$Phone)
unique_phones %>% tail(4)
```

```
## [1] "(360) 698-0417" "(206) 525-7747" "(206) 390-9205"
## [4] "(425) 557-4474"
```

```
str_detect(unique_phones, "206") %>% tail(4)
```

```
## [1] FALSE  TRUE  TRUE FALSE
```

# Replacement: `str_replace()`

What about if you want to replace a string with something else? Use `str_replace()`!

This function works very similarly to `str_detect()`, but with one extra argument:

```
str_replace(string, pattern, replacement)
```

- `replacement` is what `pattern` is substituted for.

```
str_replace(string="Hi, I'm Michael",  
            pattern="Hi", replacement="Hello")
```

```
## [1] "Hello, I'm Michael"
```

# Replacement: `str_replace()`

In the `Date` variable, let's replace each dash ("-") with an underscore ("\_")

```
dates <- restaurants$Date  
dates %>% tail(3)
```

```
## [1] "2017-03-21" "2017-03-21" "2016-10-10"
```

```
str_replace(dates, "-", "_") %>% tail(3)
```

```
## [1] "2017_03-21" "2017_03-21" "2016_10-10"
```

Wait, what?

# Replacement:

## `str_replace_all()`

`str_replace()` only changes the **first** instance of a pattern in each string!

If we want to replace **all** patterns, use `str_replace_all()`

```
dates <- restaurants$Date  
dates %>% tail(3)
```

```
## [1] "2017-03-21" "2017-03-21" "2016-10-10"
```

```
str_replace_all(dates, "-", "_") %>% tail(3)
```

```
## [1] "2017_03_21" "2017_03_21" "2016_10_10"
```

# Quick Summary

We've seen lots of functions today!

*Don't try to memorize them!* Instead, use this page as a reference.

- Character Length: `nchar` and `str_length`
- Subsetting: `substr` and `str_sub`
- Combining: `paste` and `str_c`
- Case Changes: `str_to_upper()`, `str_to_lower()`, and `str_to_title()`
- Removing Whitespace: `str_trim`
- Pattern Detection/Replacement: `str_detect()` and `str_replace()`

# Activity 1: Base R Functions

The variable `Inspection_Date` is in the format "MM/DD/YYYY". In this question, we'll change the format using functions for strings.

1. How long is each character string in this variable?
2. Use `substr()` to extract the month of each entry and save it to an object called "months"
3. Use `substr()` to extract the year of each entry and save it to an object called "years"
4. Use `paste()` to combine each month and year, separated by an underscore (`_`). Save this as a new variable in the data called "Inspection\_Date\_Formatted"

# Activity: My Answers

The variable `Inspection_Date` is in the format "MM/DD/YYYY". In this question, we'll change the format using functions for strings.

1. How long is each character string in this variable?

```
table(nchar(restaurants$Inspection_Date))
```

```
##  
##      10  
## 258000
```

2. Use `substr()` to extract the month of each entry and save it to an object called "months"

3. Use `substr()` to extract the year of each entry and save it to an object called "years"

```
months <- substr(restaurants$Inspection_Date,1,2)  
years  <- substr(restaurants$Inspection_Date,7,10)
```

# Activity: My Answers

4. Use `paste()` to combine each month and year, separated by an underscore (`_`). Save this as a new variable in the data called `"Inspection_Date_Formatted"`

```
restaurants$Inspection_Date_Formatted <-  
  paste(months, years, sep = "_")  
restaurants %>%  
  select(Name, Inspection_Date, Inspection_Date_Formatted) %>%  
  head(5)
```

```
## # A tibble: 5 × 3  
##   Name                Inspection_Date Inspection_Date_Formatted  
##   <chr>                <chr>                <chr>  
## 1 "@ THE SHACK, LLC " <NA>                NA_NA  
## 2 "10 MERCER RESTAURANT" 01/24/2017          01_2017  
## 3 "10 MERCER RESTAURANT" 01/24/2017          01_2017  
## 4 "10 MERCER RESTAURANT" 01/24/2017          01_2017  
## 5 "10 MERCER RESTAURANT" 10/10/2016          10_2016
```



# Activity 2: HW 8

Let's examine the coffee shops of King County!

1. Filter your data to only include rows in which the `Name` includes the word "coffee" (in any case!)
2. Create a new variable in your data which includes the length of the business name, after removing beginning/trailing whitespace.
3. Create a new variable in your data for the inspection year, *using a `stringr` function!*
4. Create side-by-side boxplots for the length of business name vs. year.
5. Calculate the maximum `Inspection_Score` by business and year.
6. Create a line plot of maximum score ("MaxScore") over time ("Year"), by business ("Name"). That is, you should have a single line for each business. (Don't try to label them, as there are far too many!)

# Activity: My Solutions

1. Filter your data to only include rows in which the `Name` includes the word "coffee" (in any case!)

```
coffee <- restaurants
coffee$Name <- str_to_lower(coffee$Name)
coffee <- coffee %>% filter(str_detect(Name, "coffee"))
```

# Activity: My Solutions

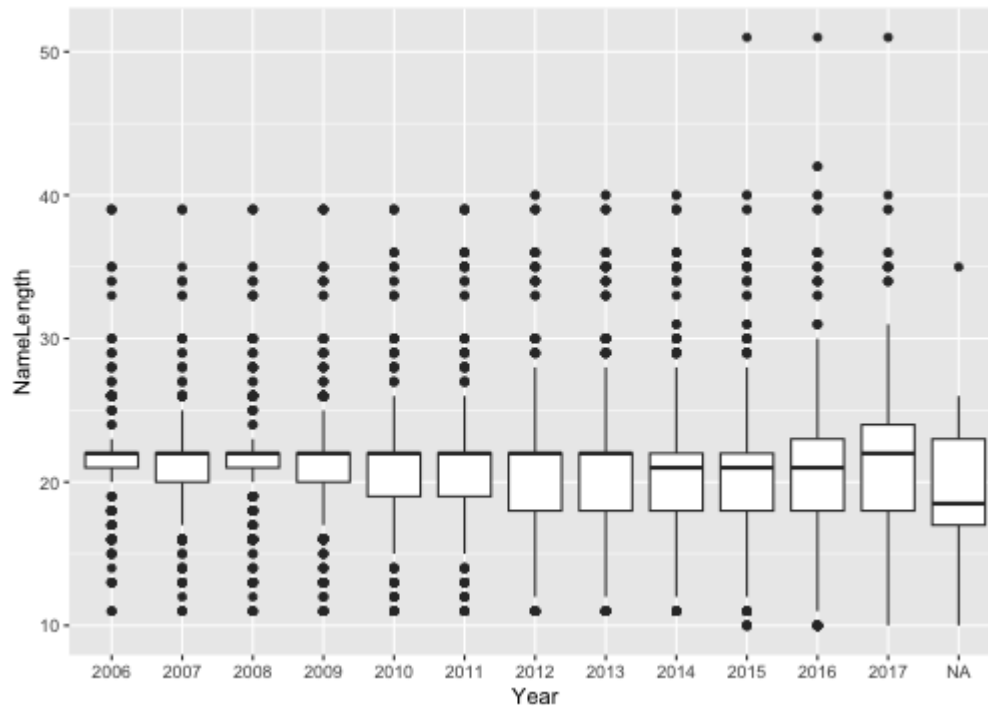
2. Create a new variable in your data which includes the length of the business name, after removing beginning/trailing whitespace.
3. Create a new variable in your data for the inspection year.

```
coffee$NameLength <- str_length(str_trim(coffee$Name))  
coffee$Year <- str_sub(coffee$Inspection_Date, -4, -1)
```

# Activity: My Solutions

4. Create side-by-side boxplots for the length of business name vs. year.

```
library(ggplot2)
ggplot(coffee, aes(Year, NameLength)) + geom_boxplot()
```



# Activity: My Solutions

5. Calculate the maximum `Inspection_Score` by business and year.

```
coffee_summary <- coffee %>% group_by(Name,Year) %>%  
  summarize(MaxScore=max(Inspection_Score))
```

```
## `summarise()` has grouped output by 'Name'. You can override using  
## the ` `.groups` argument.
```

# Activity: My Solutions

6. Create a line plot of maximum score ("MaxScore") over time ("Year"), by business ("Name"). That is, you should have a single line for each business. (Don't try to label them, as there are far too many!)

```
ggplot(coffee_summary,aes(Year,MaxScore,group=Name))+  
  geom_line(alpha=.2)
```

